

Software 2 (Theory)

Lecture 10: B-trees

Today's Topics

- Discuss structure of B-Trees
- The properties of B-Trees
- Basic operations on B-Trees
 - Searching, Insertion and Deletion
- Understand some applications of B-Trees

Introduction

- B-trees are powerful data structures used in various applications, including databases and file systems.
- A B-tree is a self-balancing tree data structure that maintains sorted data and facilitates efficient searching, insertion, and deletion operations.
- The "B" in B-tree can stand for "balanced," "binary," or "Bayer," after its inventors Rudolf Bayer and Edward M. McCreight.
- Unlike binary search trees, B-trees can have multiple keys and child pointers per node, leading to a more balanced structure.

Structure of B-Tree

- Each node in a B-tree can have multiple keys and child pointers.
 - Also known as “large key” trees.
- The keys within a node are stored in sorted order, facilitating efficient searching.
 - Offers larger branching factor and shallower height for faster search and insertion.
- All leaf nodes are at the same level.
 - ensuring a balanced tree structure.
- The number of keys within a node is constrained by properties such as the minimum and maximum number of keys allowed per node.
 - Maintaining a balance which guarantees the time complexity for its operations.

Time Complexity of B-Tree

- **Searching:** Search for a key by traversing the tree from the root downwards, following the appropriate child pointers based on the key values.
- **Insertion:** Insert a new key into the B-tree by finding the appropriate leaf node and inserting the key while maintaining the sorted order.
- **Deletion:** Delete a key from the B-tree by first locating the key and then removing it from the appropriate node, adjusting the structure if necessary to maintain balance.

Algorithm	Time Complexity
Search	$O(\log n)$
Insert	$O(\log n)$
Delete	$O(\log n)$

Properties of B-Trees

- **Balanced Structure:** B-trees maintain a balanced structure by redistributing data among nodes as necessary, ensuring efficient operations.
- **Efficient Searching:** Searching in a B-tree is efficient, typically requiring logarithmic time complexity due to the balanced nature of the tree.
- **Efficient Insertion and Deletion:** B-trees support efficient insertion and deletion operations by dynamically adjusting the structure while maintaining balance.
- **Optimal Disk I/O:** B-trees are well-suited for applications requiring persistent storage on disk, as they minimise disk I/O by ensuring a balanced tree structure.

B-Trees shape properties

A B-tree of order m is defined to have the following shape properties:


- Every node has at most m children.
- Every non-leaf node (except the root node) has at least $m = 2$ children.
- The root node, if it is not a leaf node, has at least two children.
- A non-leaf node with m children contains $m - 1$ search keys.
- Each internal node, except for the root, has between $\lceil m/2 \rceil$ and m children.
- All leaves are at the same level in the tree, carries information and hence, the tree is always height balanced.

B-Tree Example with order 5 ($m = 5$)

The root has between 2 and m children.


Each non-root internal node has between $\lceil m/2 \rceil$ and m children.

All external nodes are at the same level. (External nodes are represented by null pointers in implementations.)




B-Tree Insertion : insert 10


- We find the location for 10 by following a path from the root using the stored key values to guide the search.
- The search falls out the tree at the 4th child of the 1st child of the root.
- The 1st child of the root has room for the new element, so we store it there.



B-Tree Insertion : insert 11


- We fall out of the tree at the child to the right of key 10.
- But there is no more room in the left child of the root to hold 11.
- Therefore, we must split this node.



B-Tree Insertion : insert 11 (continued)


The $m + 1$ children are divided evenly between the old and new nodes.

The parent gets one new child. (If the parent become overfull, then it, too, will have to be split).

B-Tree deletion : delete 8



Removing 8 might force us to move another key up from one of the children. It could either be the 3 from the 1st child or the 10 from the second child.

However, neither child has more than the minimum number of children ($m = \lceil m/2 \rceil = 3$) for $m = 5$, so the two nodes will have to be merged. Nothing moves up.





B-Tree Deletion : delete 13

- Deleting 13 would cause the node containing it to become underfill.
- To fix this, we try to reassign one key from a sibling that has spares.
- The 13 is replaced by the parent's key 12.
- The parent's key 12 is replaced by the spare key 11 from the left sibling.




B-Tree Deletion : delete 11

- 11 is in a non-leaf, so replace it by the value immediately preceding: 10.
- 10 is at leaf, and this node has spares.





B-Tree Deletion : delete 2

- Although 2 is at leaf level, removing it leads to an underfill node.
- The node has no left sibling. It does have a right sibling, but that node is at its minimum occupancy already.
- Therefore, the node must be merged with its right sibling.



B-Tree Deletion : delete 2

- This is no longer a B-tree, because the root does not have at least 2 children.
- Therefore, we must remove the root, making its child the new root.



Applications of B-Trees

- B-trees are widely used in databases and file systems due to their efficient searching, insertion, and deletion operations.
- They are also used in other applications such as indexing and caching, where fast access to large amounts of data is required.
- By maintaining a balanced structure and supporting efficient operations, B-trees enable high-performance applications in various domains.
- B-Trees are used in CAD systems to organize and search geometric data.
- B-Trees are also used in other areas such as natural language processing, computer networks, and cryptography.

Summary

- Understand:
 - B-Trees
 - B-Tree search, insertion and deletion
 - Similar time complex for the 3 operations
 - Some applications of B-Trees
- Read
 - Skiena, Section 15.1
 - Goodrick, Chapter 15
 - Cormen, Chapter 18
- Next
 - Revision Lecture